

Principles of Database Systems (CS307)

Lecture 14: Query Optimization

Zhong-Qiu Wang

Department of Computer Science and Engineering
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.
- The slides are largely based on the slides provided by Dr. Yuxin Ma

Announcements

- Assignment on Trigger, due date: 23:59pm on December 15th 2024, Beijing Time
- For Project II, if you choose to submit
 - by 23:30pm on December 22nd 2024 (Sunday), Beijing time
 - Your project defense has to be in your lab session that is in the same week as the above deadline
 - Your score will be multiplied by 1.1
 - by 23:30pm on December 29th 2024 (Sunday), Beijing time
 - Your project defense has to be in your lab session that is in the same week as the above deadline
 - Your score will be multiplied by 1.0

Query Optimization

- Purpose of query optimization
 - Select an effective way to retrieve the data based on queries while spending the least computational effort
 - However, it is only “spending less computational effort” in most scenarios, not least

Query Optimization

- Users don't need to consider the best way of writing queries
 - We want DBMS to construct a query-evaluation plan that minimizes the cost of query evaluation
- Automated optimization can perform better (for most of the time)
 - Utilize the data dictionary
 - Real-time utilization based on physical storage changes
 - Optimizer can evaluate hundreds of execution plans in a very short time compared with human programmers
 - Human users do not need to learn advanced optimization techniques any more, which is conducted by optimizers instead

An Example in the Movie Dataset

- The same query can be represented in different plans
 - E.g., retrieve the titles of those movies from China



```
select m.title  
from movies m, countries c  
where m.country = c.country_code and c.country_name = 'China';
```

An Example in the Movie Dataset

- The corresponding relational algebra expressions:

(1) $\Pi_{title} (\sigma_{movies.country = countries.country_code \wedge countries.country = "China"}(movies \times countries))$

(2) $\Pi_{title} (\sigma_{countries.country = "China"}(movies \bowtie_{movies.country = countries.country_code} countries))$

(3) $\Pi_{title} (movies \bowtie_{movies.country = countries.country_code} (\sigma_{countries.country = "China"}(countries)))$

An Example in the Movie Dataset

- The corresponding relational algebra expressions:

(1) $\Pi_{title} (\sigma_{movies.country = countries.country_code \wedge countries.country = "China"}(movies \times countries))$

(2) $\Pi_{title} (\sigma_{countries.country = "China"}(movies \bowtie_{movies.country = countries.country_code} countries))$

(3) $\Pi_{title} (movies \bowtie_{movies.country = countries.country_code} (\sigma_{countries.country = "China"}(countries)))$

- In (1), a **full Cartesian product** will be computed, which costs huge time for matching all pairs and massive temporary storage space for the intermediate product table
- In (2), a **smaller intermediate join table** is to be cached, which saves some space
- In (3), **the filter** $(\sigma_{c.country = "China"})$ **reduces the size of the right table** in the join operation, which saves a lot of time for pair matching and caching intermediate join table

An Example in the Movie Dataset

- In addition, the filter operation can be further accelerated once an index is built upon the *country* column

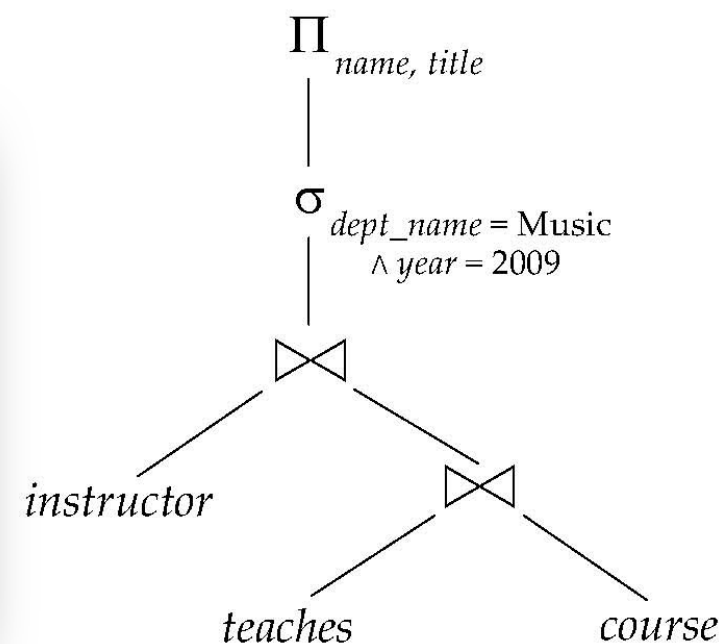
(3) $\Pi_{title} (\text{movies} \bowtie_{\text{movies.country} = \text{countries.country_code}} (\sigma_{\text{countries.country} = \text{"China"}} (\text{countries})))$

Generating Equivalent Expressions

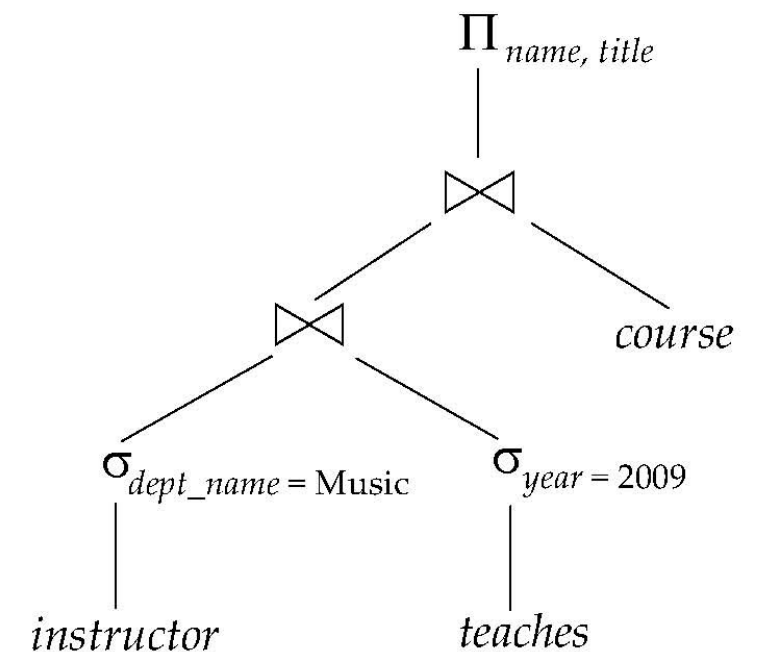
- Alternative ways of evaluating a given query
 - Equivalent expressions
 - Different algorithms for each operation



```
select name, title
from instructor
  natural join (teaches natural join course)
where dept_name = 'Music' and year = 2009;
```



(a) Initial expression tree



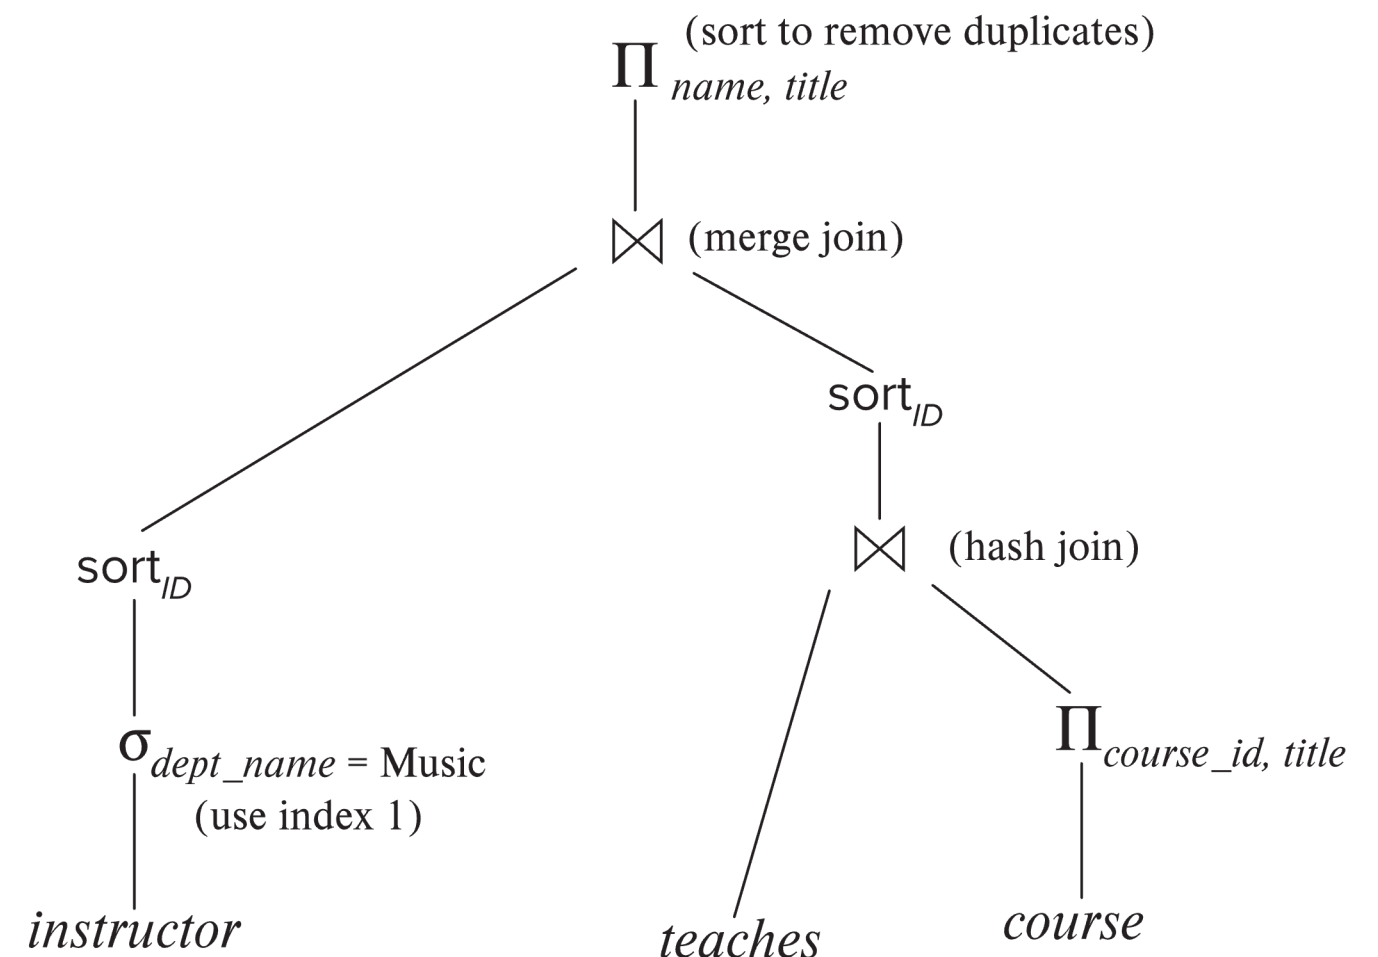
(b) Tree after multiple transformations

Generating Equivalent Expressions

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated



```
select name, title
from instructor
      natural join (teaches natural join course)
where dept_name = 'Music' and year = 2009;
```



Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
 - Note: order of tuples is irrelevant
 - We don't care if they generate different results on databases that violate integrity constraints
- An **equivalence rule** says that expressions of two forms are equivalent
 - ... i.e., we can replace expression of the first form by second, or vice versa
 - The optimizer uses equivalence rules to transform expressions into other logically equivalent expressions

Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are **commutative**

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) \equiv \Pi_{L_1}(E)$$

where $L_1 \subseteq L_2 \dots \subseteq L_n$

4. Selections can be combined with Cartesian products and theta joins

a) $\sigma_{\theta}(E_1 \times E_2) \equiv E_1 \bowtie_{\theta} E_2$ (same as the definition of theta-join)

b) $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) \equiv E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

Equivalence Rules

5. Theta-join operations (and natural joins) are commutative.

$$E_1 \bowtie E_2 \equiv E_2 \bowtie E_1$$

6. (a) Natural join operations are **associative**:

$$(E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$$


(b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 \equiv E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$



where θ_2 involves attributes from only E_2 and E_3

Equivalence Rules



- Representation of Rule 5, 6(a) and 6(b) with diagrams



Rule 5




Rule 6.a



Rule 7.a

If θ only has attributes from E_1



Equivalence Rules

7. The selection operation distributes over the theta join operation under the following two conditions:

- (a) When all the attributes in θ_0 involve only the attributes of one of the expressions (E_1) being joined:

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

- (b) When θ_1 involves only the attributes of E_1 and θ_2 involves only the attributes of E_2 :

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

Equivalence Rules

8. The projection operation distributes over the theta join operation as follows:

(a) If θ involves only attributes from $L_1 \cup L_2$:

- Let L_1 and L_2 be sets of attributes from E_1 and E_2 , respectively,

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1}(E_1) \bowtie_{\theta} \Pi_{L_2}(E_2)$$

(b) In general, consider a join $E_1 \bowtie_{\theta} E_2$:

- Let L_1 and L_2 be sets of attributes from E_1 and E_2 , respectively,
- Let L_3 be attributes of E_1 that are involved in join condition θ , but are not in L_1 , and,
- Let L_4 be attributes of E_2 that are involved in join condition θ , but are not in L_2 :

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1 \cup L_2}(\Pi_{L_1 \cup L_3}(E_1) \bowtie_{\theta} \Pi_{L_2 \cup L_4}(E_2))$$

* Similar equivalences hold for left, right, and full outer join operations: \bowtie , \ltimes , and \rhd

Equivalence Rules

9. The set operations union and intersection are commutative

$$E_1 \cup E_2 \equiv E_2 \cup E_1$$

$$E_1 \cap E_2 \equiv E_2 \cap E_1$$

- However, set difference is not commutative

10. Set union and intersection are associative

$$(E_1 \cup E_2) \cup E_3 \equiv E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 \equiv E_1 \cap (E_2 \cap E_3)$$

Equivalence Rules

11. The selection operation distributes over \cup , \cap and $-$

(a) $\sigma_{\theta} (E_1 \cup E_2) \equiv \sigma_{\theta} (E_1) \cup \sigma_{\theta}(E_2)$

(b) $\sigma_{\theta} (E_1 \cap E_2) \equiv \sigma_{\theta} (E_1) \cap \sigma_{\theta}(E_2)$

(c) $\sigma_{\theta} (E_1 - E_2) \equiv \sigma_{\theta} (E_1) - \sigma_{\theta}(E_2)$

(d) $\sigma_{\theta} (E_1 \cap E_2) \equiv \sigma_{\theta}(E_1) \cap E_2$

(e) $\sigma_{\theta} (E_1 - E_2) \equiv \sigma_{\theta}(E_1) - E_2$

12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) \equiv (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$

Transformation Example: Pushing Selections

- Query: Find the names of all instructors in the Music department, along with the titles of the courses (in the Music department) that they teach

```
select name, title
from instructor natural join (teaches natural join course
where dept_name = 'Music');
```

- $\Pi_{name, title}(\sigma_{dept_name = 'Music'}(instructor \bowtie (teaches \bowtie \Pi_{course_id, title}(course))))$
- Transformation using rule 7(a):
 - $\Pi_{name, title}((\sigma_{dept_name = 'Music'}(instructor)) \bowtie (teaches \bowtie \Pi_{course_id, title}(course)))$

Perform selection as early as possible reduces the size of the relation to be joined

Transformation Example: Multiple Transformations

- Query: Find the names of all instructors in the Music department who have taught a course in 2017, along with the titles of the courses that they taught

- $\Pi_{name, title}(\sigma_{dept_name = \text{"Music"} \wedge year = 2017} (instructor \bowtie (teaches \bowtie \Pi_{course_id, title} (course))))$

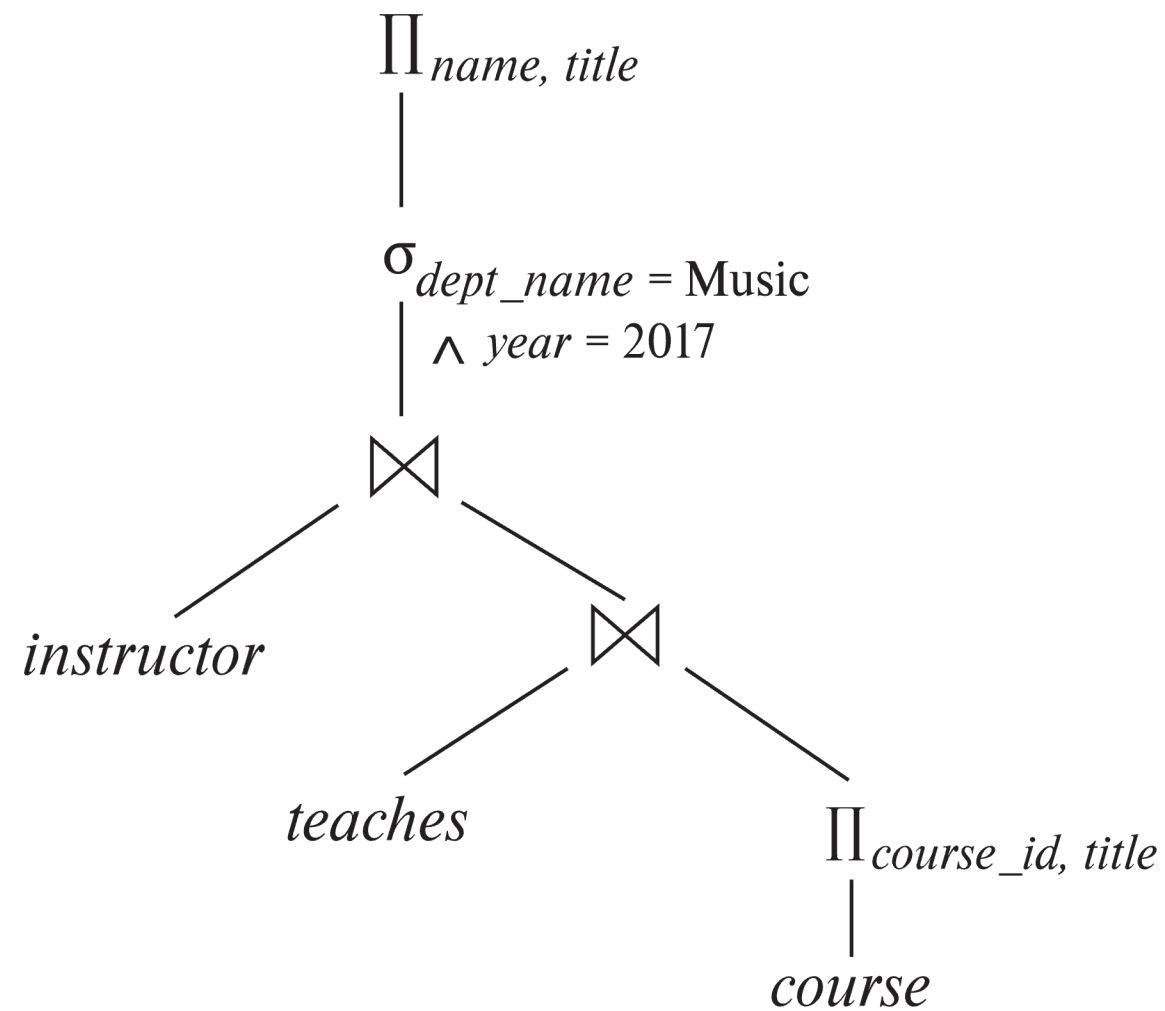
- Transformation using join associatively (Rule 6(a)):

- $\Pi_{name, title}(\sigma_{dept_name = \text{"Music"} \wedge year = 2017} ((instructor \bowtie teaches) \bowtie \Pi_{course_id, title} (course)))$

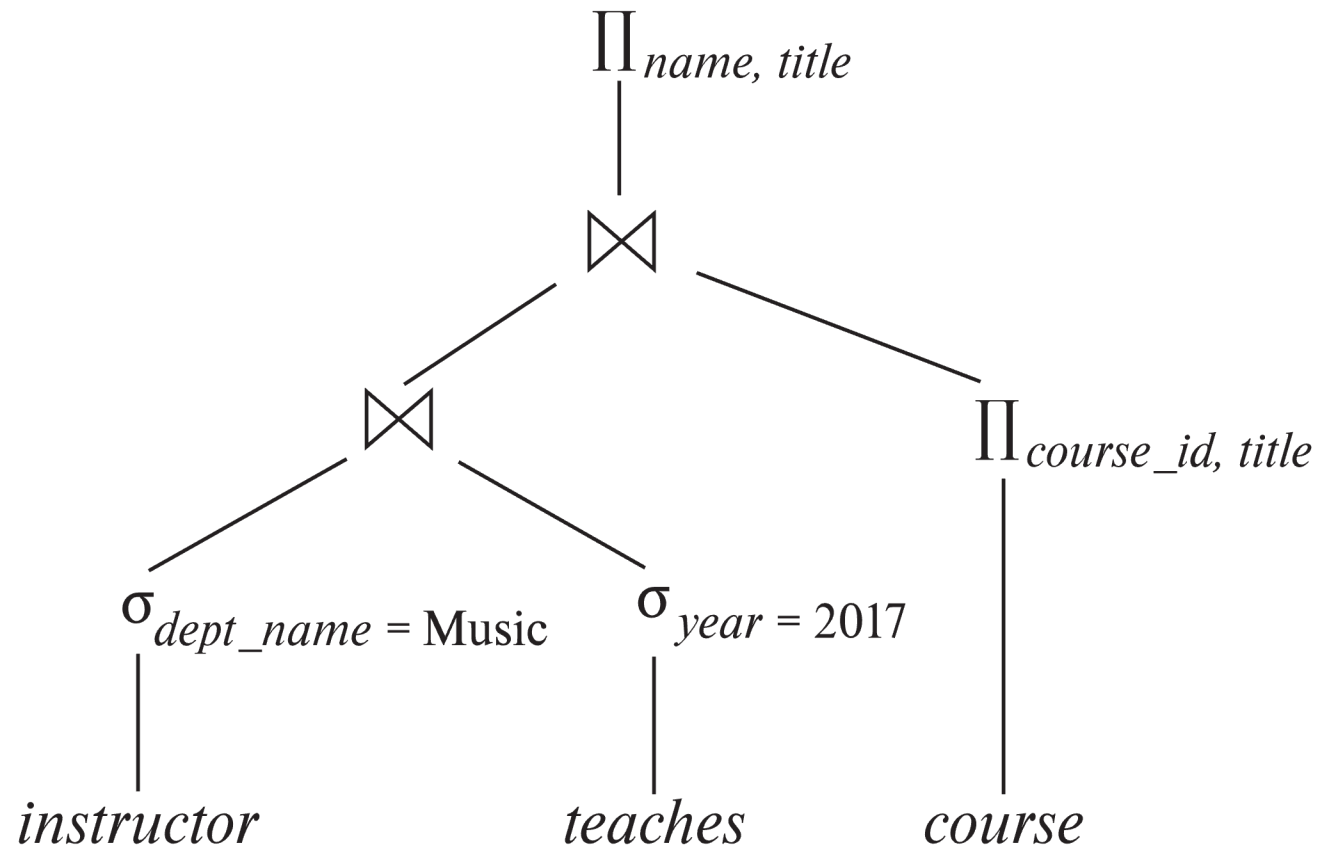
- Second form provides an opportunity to apply the “perform selections early” rule, resulting in the subexpression:

- $\sigma_{dept_name = \text{"Music"}} (instructor) \bowtie \sigma_{year = 2017} (teaches)$

Transformation Example: Multiple Transformations



(a) Initial expression tree



(b) Tree after multiple transformations

* Transformation Example: Pushing Projections

- Consider $\Pi_{name, title} (\sigma_{dept_name = \text{"Music"}} (instructor) \bowtie teaches) \bowtie \Pi_{course_id, title} (course))$

- When we compute

$$(\sigma_{dept_name = \text{"Music"}} (instructor) \bowtie teaches),$$

we obtain a relation whose schema is:

$(ID, name, dept_name, salary, course_id, sec_id, semester, year)$

- Push projections using equivalence rules 8a and 8b; eliminate unneeded attributes from intermediate results to get:

$$\Pi_{name, title} (\Pi_{name, course_id} (\sigma_{dept_name = \text{"Music"}} (instructor) \bowtie teaches) \bowtie \Pi_{course_id, title} (course)))$$

Perform projections as early as possible reduces the size of the relation to be joined

Join Ordering Example

- For all relations r_1, r_2 , and r_3 ,

$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

- * (Join Associativity) \bowtie
- If $r_2 \bowtie r_3$ is quite large and $r_1 \bowtie r_2$ is small, we choose
$$(r_1 \bowtie r_2) \bowtie r_3$$
so that we compute and store a smaller temporary relation

Cost Estimation

- Cost difference between evaluation plans for a query can be enormous
 - E.g., seconds vs. days in some cases
- Steps in cost-based query optimization
 - 1. Generate logically equivalent expressions using equivalence rules
 - 2. Annotate resultant expressions to get alternative query plans
 - 3. Choose the cheapest plan based on estimated cost

Cost Estimation

- Estimation of plan cost based on:
 - Statistical information about relations, such as:
 - number of tuples, number of distinct values for an attribute
 - Statistics estimation for intermediate results
 - to compute cost of complex expressions
 - Cost formulae for algorithms, computed using statistics

For more, please refer to Section 16.3 “Estimating Statistics of Expression Results” in the reference textbook